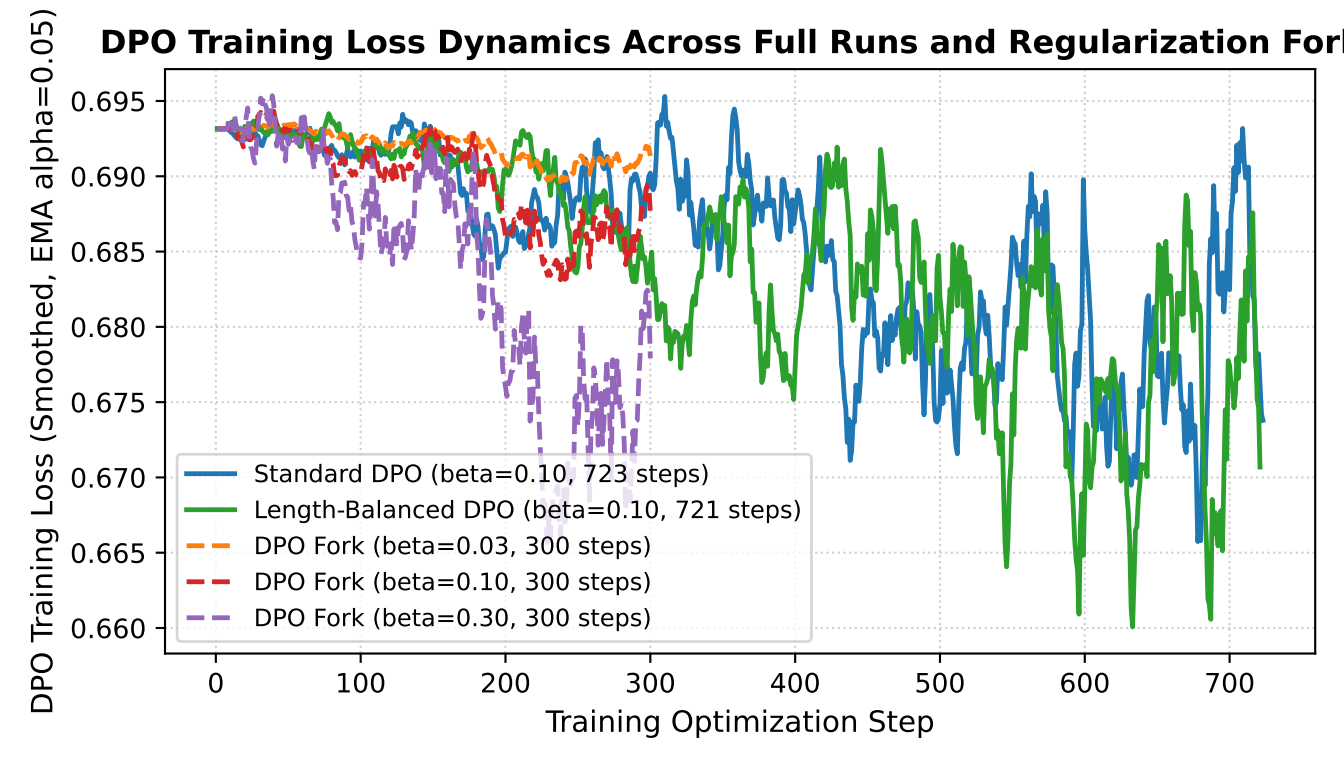
DPO Training Loss Dynamics Across Full Runs and Regularization For
DPO Training Loss (Smoothed, EMA alpha=0.05)
0.695
0.690
0.685
0.680
0.675
0.670
0.665
0.660
0
100
200
300
400
500
600
700
Training Optimization Step
Standard DPO (beta=0.10, 723 steps)
Length-Balanced DPO (beta=0.10, 721 steps)
DPO Fork (beta=0.03, 300 steps)
DPO Fork (beta=0.10, 300 steps)
DPO Fork (beta=0.30, 300 steps)