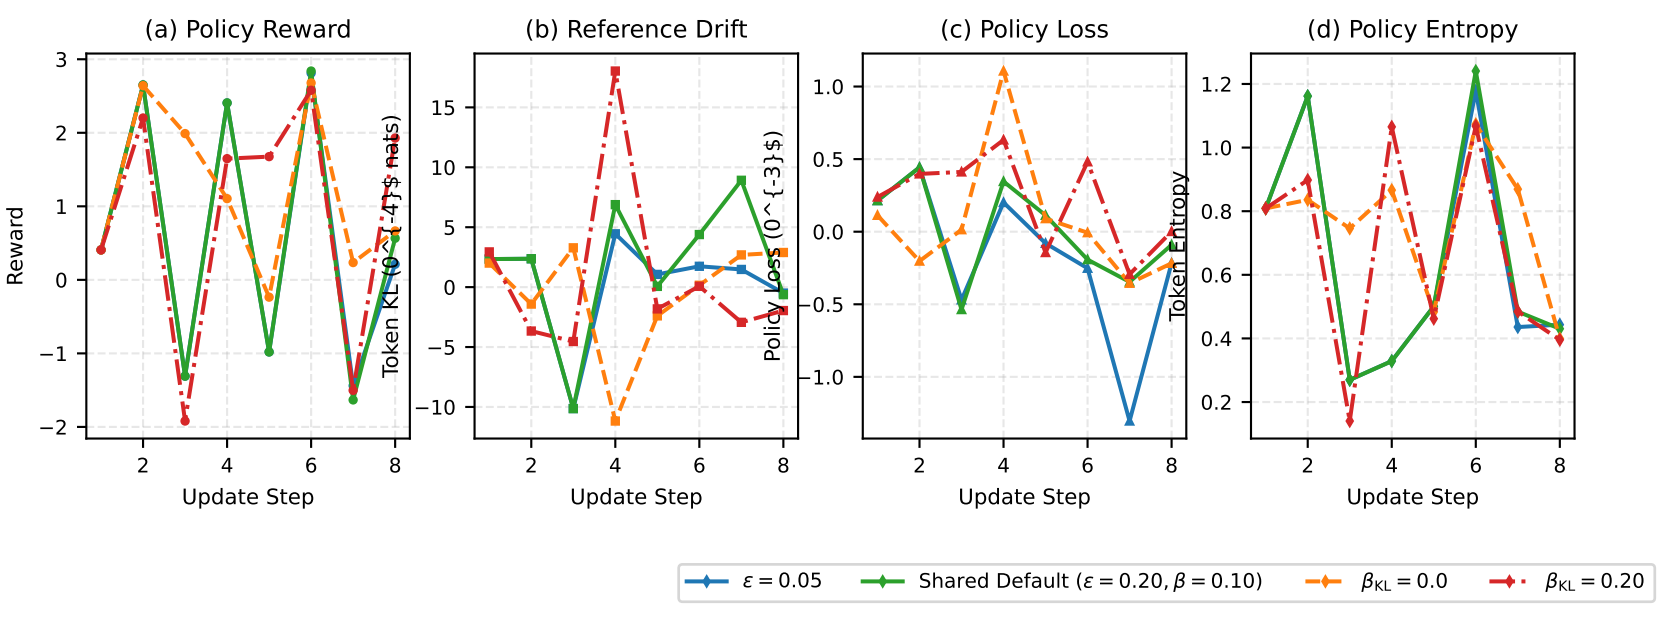
(a) Policy Reward
Reward
Update Step
(b) Reference Drift
Token KL (10^{-4} nats)
Update Step
(c) Policy Loss
Policy Loss (10^{-3})
Update Step
(d) Policy Entropy
Token Entropy
Update Step
$\epsilon = 0.05$
Shared Default ($\epsilon = 0.20, \beta = 0.10$)
$\beta_{KL} = 0.0$
$\beta_{KL} = 0.20$