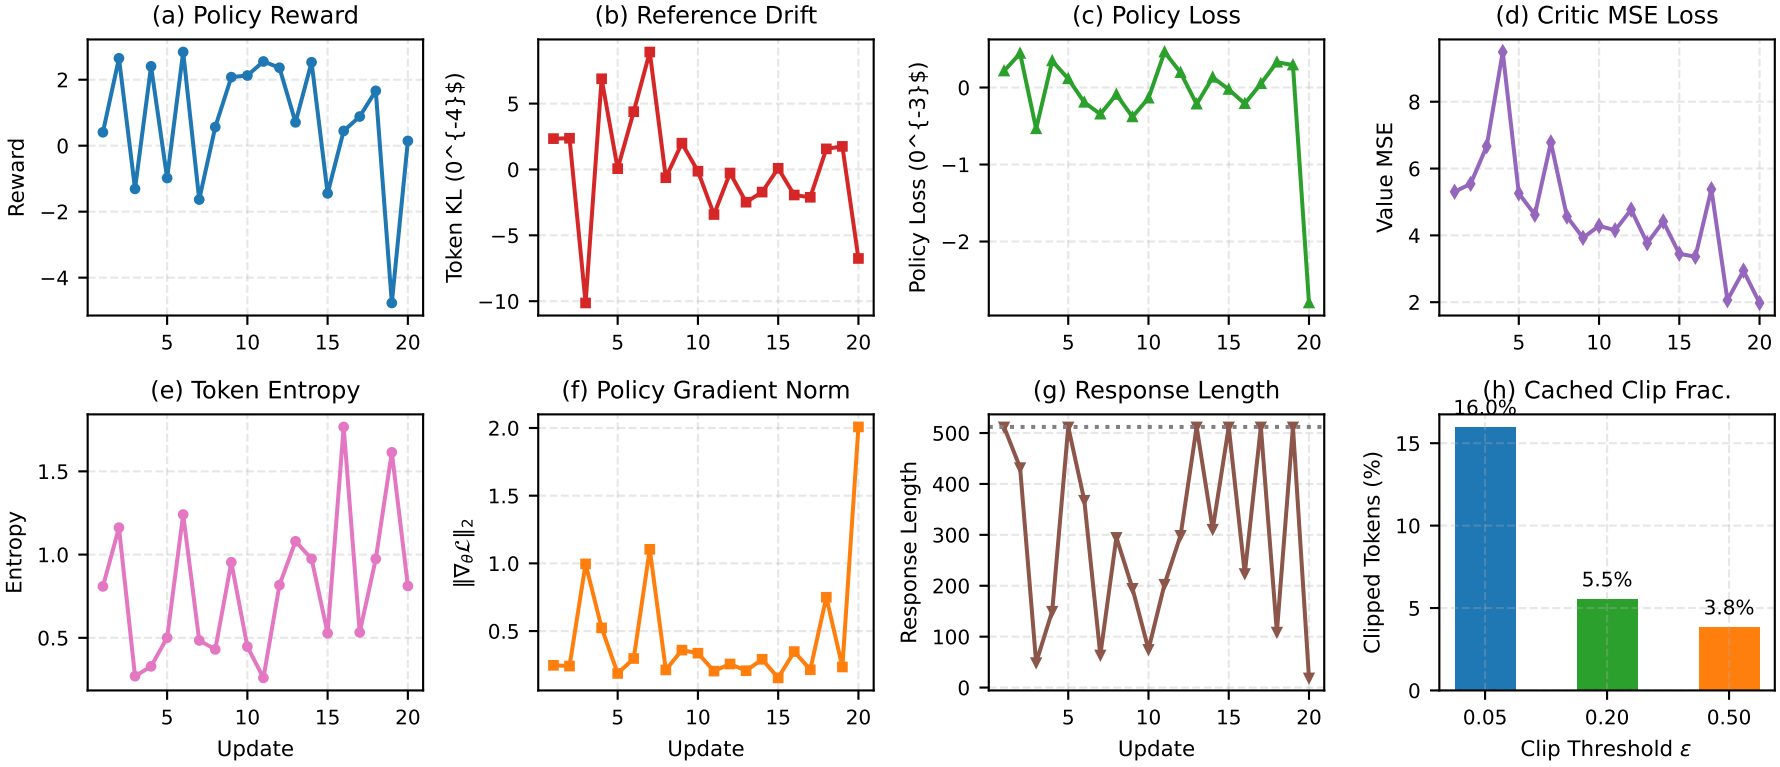
(a) Policy Reward
Reward
(b) Reference Drift
Token KL (0^{-4}$)
(c) Policy Loss
Policy Loss (0^{-3}$)
(d) Critic MSE Loss
Value MSE
(e) Token Entropy
Entropy
Update
(f) Policy Gradient Norm
$\|\nabla_\theta \mathcal{L}\|_2$
Update
(g) Response Length
Response Length
Update
(h) Cached Clip Frac.
Clipped Tokens (%)
16.0%
5.5%
3.8%
0.05
0.20
0.50
Clip Threshold ϵ